



Operating Systems Lab

CSCE 000/3402

Lab Lecture 2: The Fork System Call Implementing a Shell


Karim Sobh

kmsobh@aucegypt.edu



Multiprogramming

- The OS is loaded into memory simultaneously with other jobs.
- **Multiplexed resources among jobs:**
 - A job yields the CPU upon I/O.
 - OS assigns another job to CPU.
 - Next job needs to be already loaded in memory.
- **Protections techniques:**
 - Protect jobs from each other.
 - Protect the OS from other jobs.
- Transparent job submission and results collection.
- **Non-preemptive** → not interactive.
- Systems are referred to as **minicomputers**.



- **Definition of concurrency:** jobs are considered concurrent by being loaded simultaneously in memory.

Note: in a uniprocessor environment only one job can have the CPU at any point in time



Job Scheduling

- Deciding on which job runs next.
- Different scheduling algorithms are available.
- Scheduling types:
 - **Non-preemptive Scheduling:**
 - The job leaves the CPU only at its own will.
 - No mechanism for the OS to kick out a running job from the CPU forcefully.
 - **Preemptive Scheduling:**
 - CPU time can be sliced and shared evenly among different jobs.
 - A mechanism for giving the OS control to kick out a running job from the CPU.
 - Needs hardware support; timers and interrupts.
 - Introduced the notion of a **Timeshare** system.



Timesharing

- Timesharing is a variant of multiprogramming.
- Needed for quick response time.
- Based on users with terminals trying to access shared resources.
- Built for interactive processing based on tasks that can be split into short duration transactions.
- Basically I/O was introduced in the equation.
- The seed for utility computing; like the electrical grid.



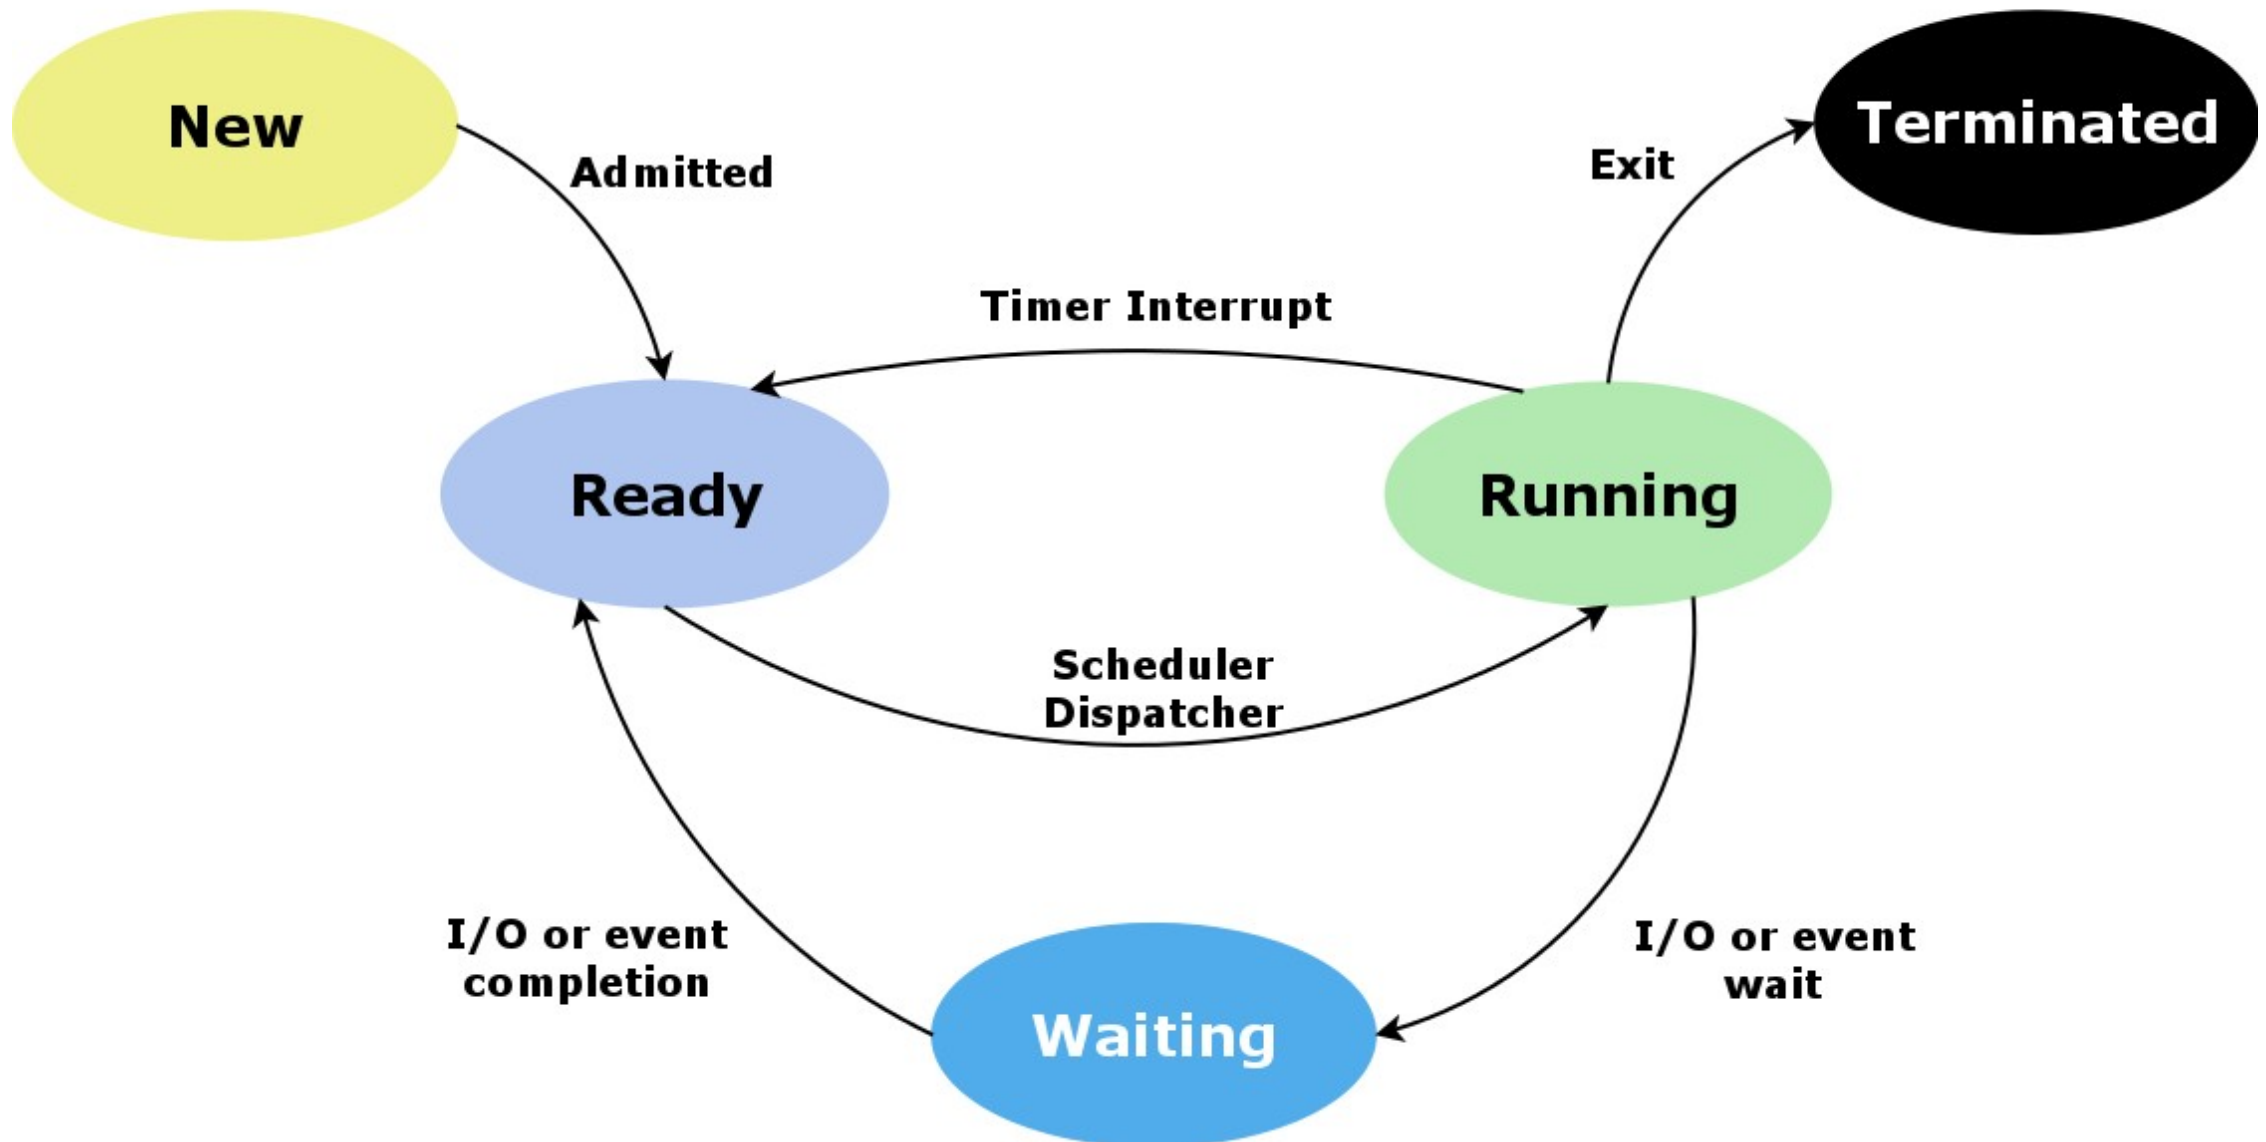
Process Creation

- Four principle events cause processes to be created:
 1. System initialization.
 2. Execution of a process-creation system call by a running process.
 3. A user request to create a new process.
 4. Initiation of a batch job; applies to main frame when different users submit batch jobs and the OS starts one after the other.
- Usually a global parent process is created and the rest of processes are its children forming a hierarchy; e.g. UNIX init process started at boot time.
- The mechanism of creating the process address space differs from one OS to another; e.g. **Copy-on-Write** → **more on this later.**



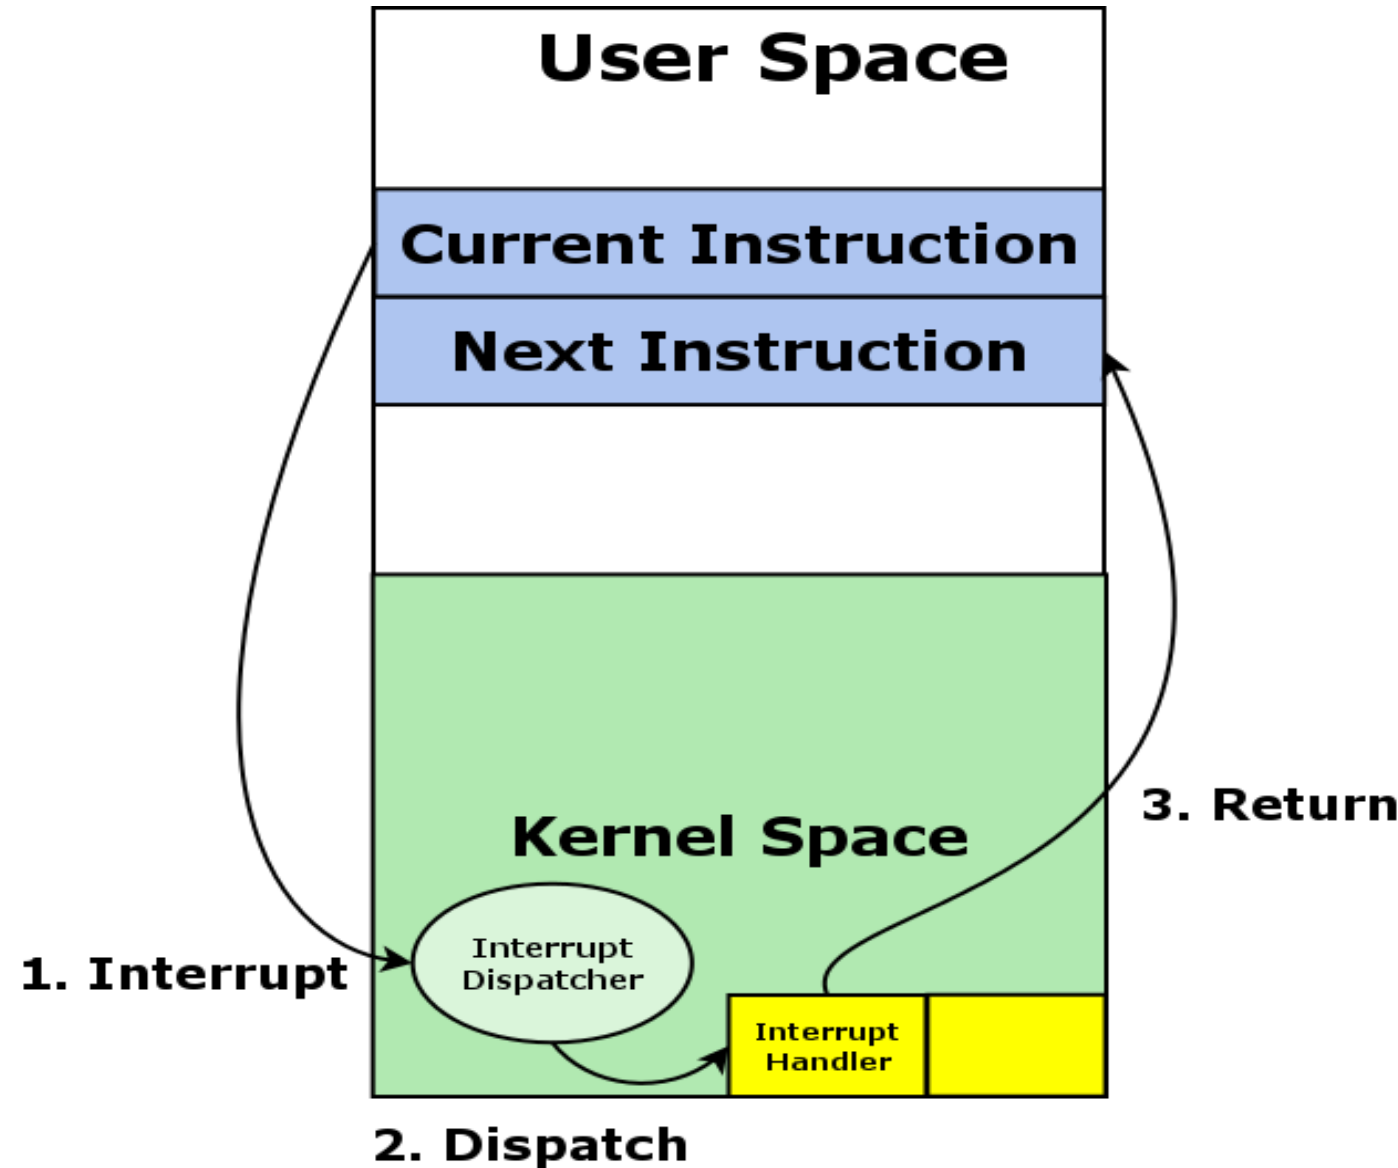
Process States

Process State




I/O Interrupts

- A program running in the user space requests I/O.
- A software interrupt is generated through an invocation of a system call.
- Execution mode is switched to kernel space.
- Kernel dispatch the interrupt to the interrupt handler.
- After servicing the interrupt execution mode is switched back to user mode.
- Continue from the next instruction in the user program after the instruction that initiated the I/O.



Pipes

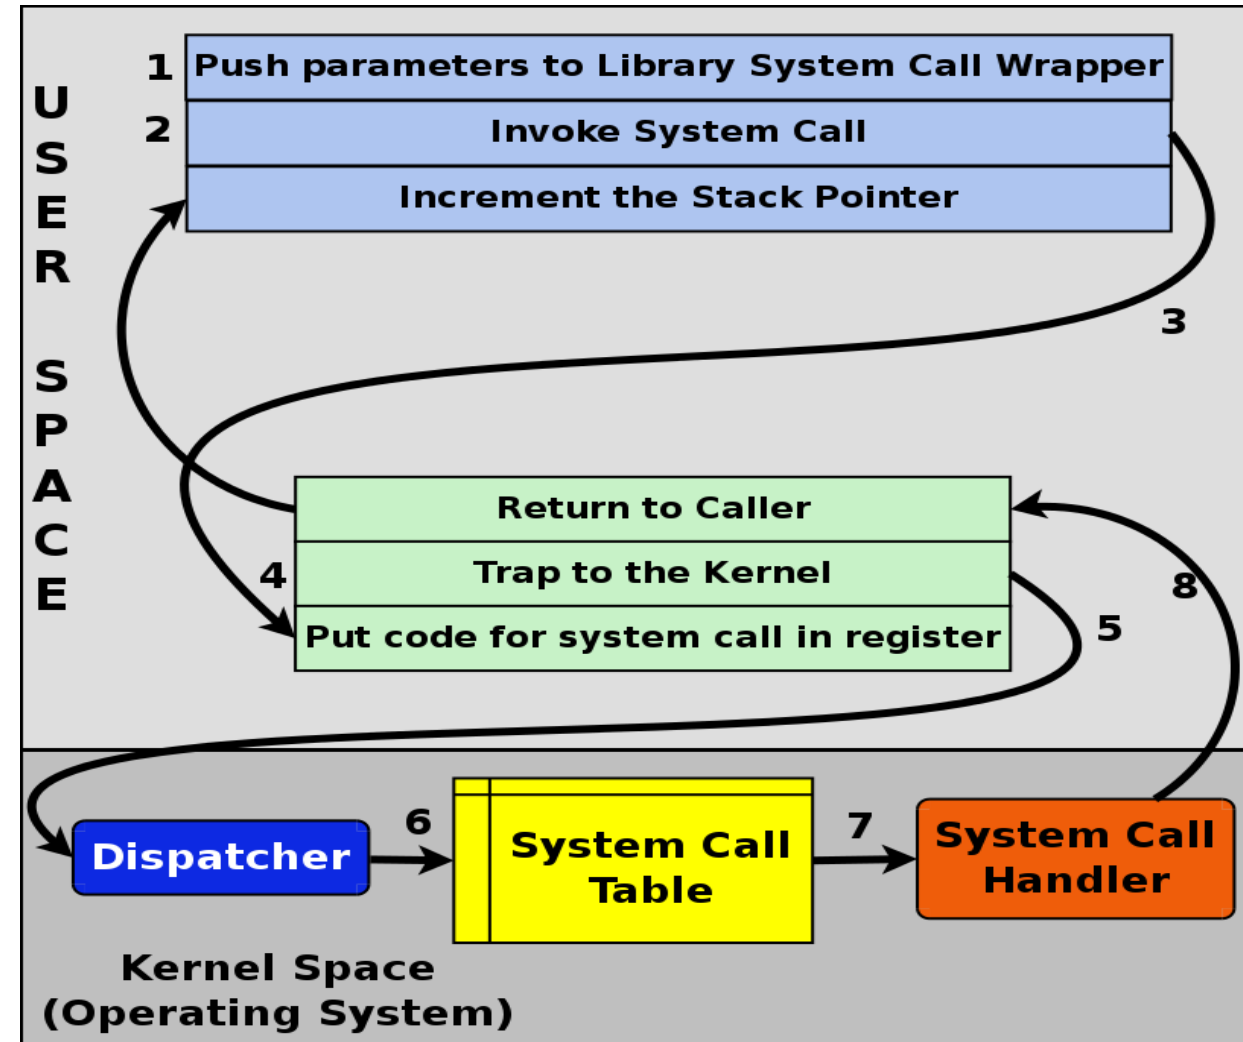


- A mechanism for interprocess communication over files.
- Pipes are a sort of pseudo-files that can connect two processes.
- Processes can exchange data/messages over pipes.
- Pipes need to be set in advance.



System Calls

- A well defined **entry point** for user programs **to kernel mode** and allows invoking kernel routines.
- A system call is basically a **function wrapper**.
- Implemented as user mode **library routine**.
- User programs invoke library routines and pass parameters to them.
- A library routine initiates a **software interrupt**; trap.
- Pass arguments to the kernel.
- Kernel execute privileged task.
- Return to the Library and consequently to user program.



System Call Examples

- Process Management:
 - `pid = fork()`
 - `pid = waitpid(pid,&statloc,options)`
 - `s = execve(name,argv,environp)`
 - `exit(status)`
- File Management:
 - `fd = open(file,how,...)`
 - `s = close(fd)`
 - `n = read (fd,buffer,nbytes)`
 - `N = write (fd,buffer,nbytes)`
 - `position = lseek (fd,offset,whence)`
 - `s = stat(name,&buf)`
- Directory and Filesystem management:
 - `s = mkdir(name,mode)`
 - `s = rmdir(name)`
 - `s = link(name1,name2)`
 - `s = unlink(name)`
 - `s = mount(name,flag)`
 - `s = umount(name)`
- Miscellaneous:
 - `s = chdir(dirname)`
 - `s = chmod(name,mode)`
 - `s = kill(pid,signal)`
 - `seconds = time(&seconds)`

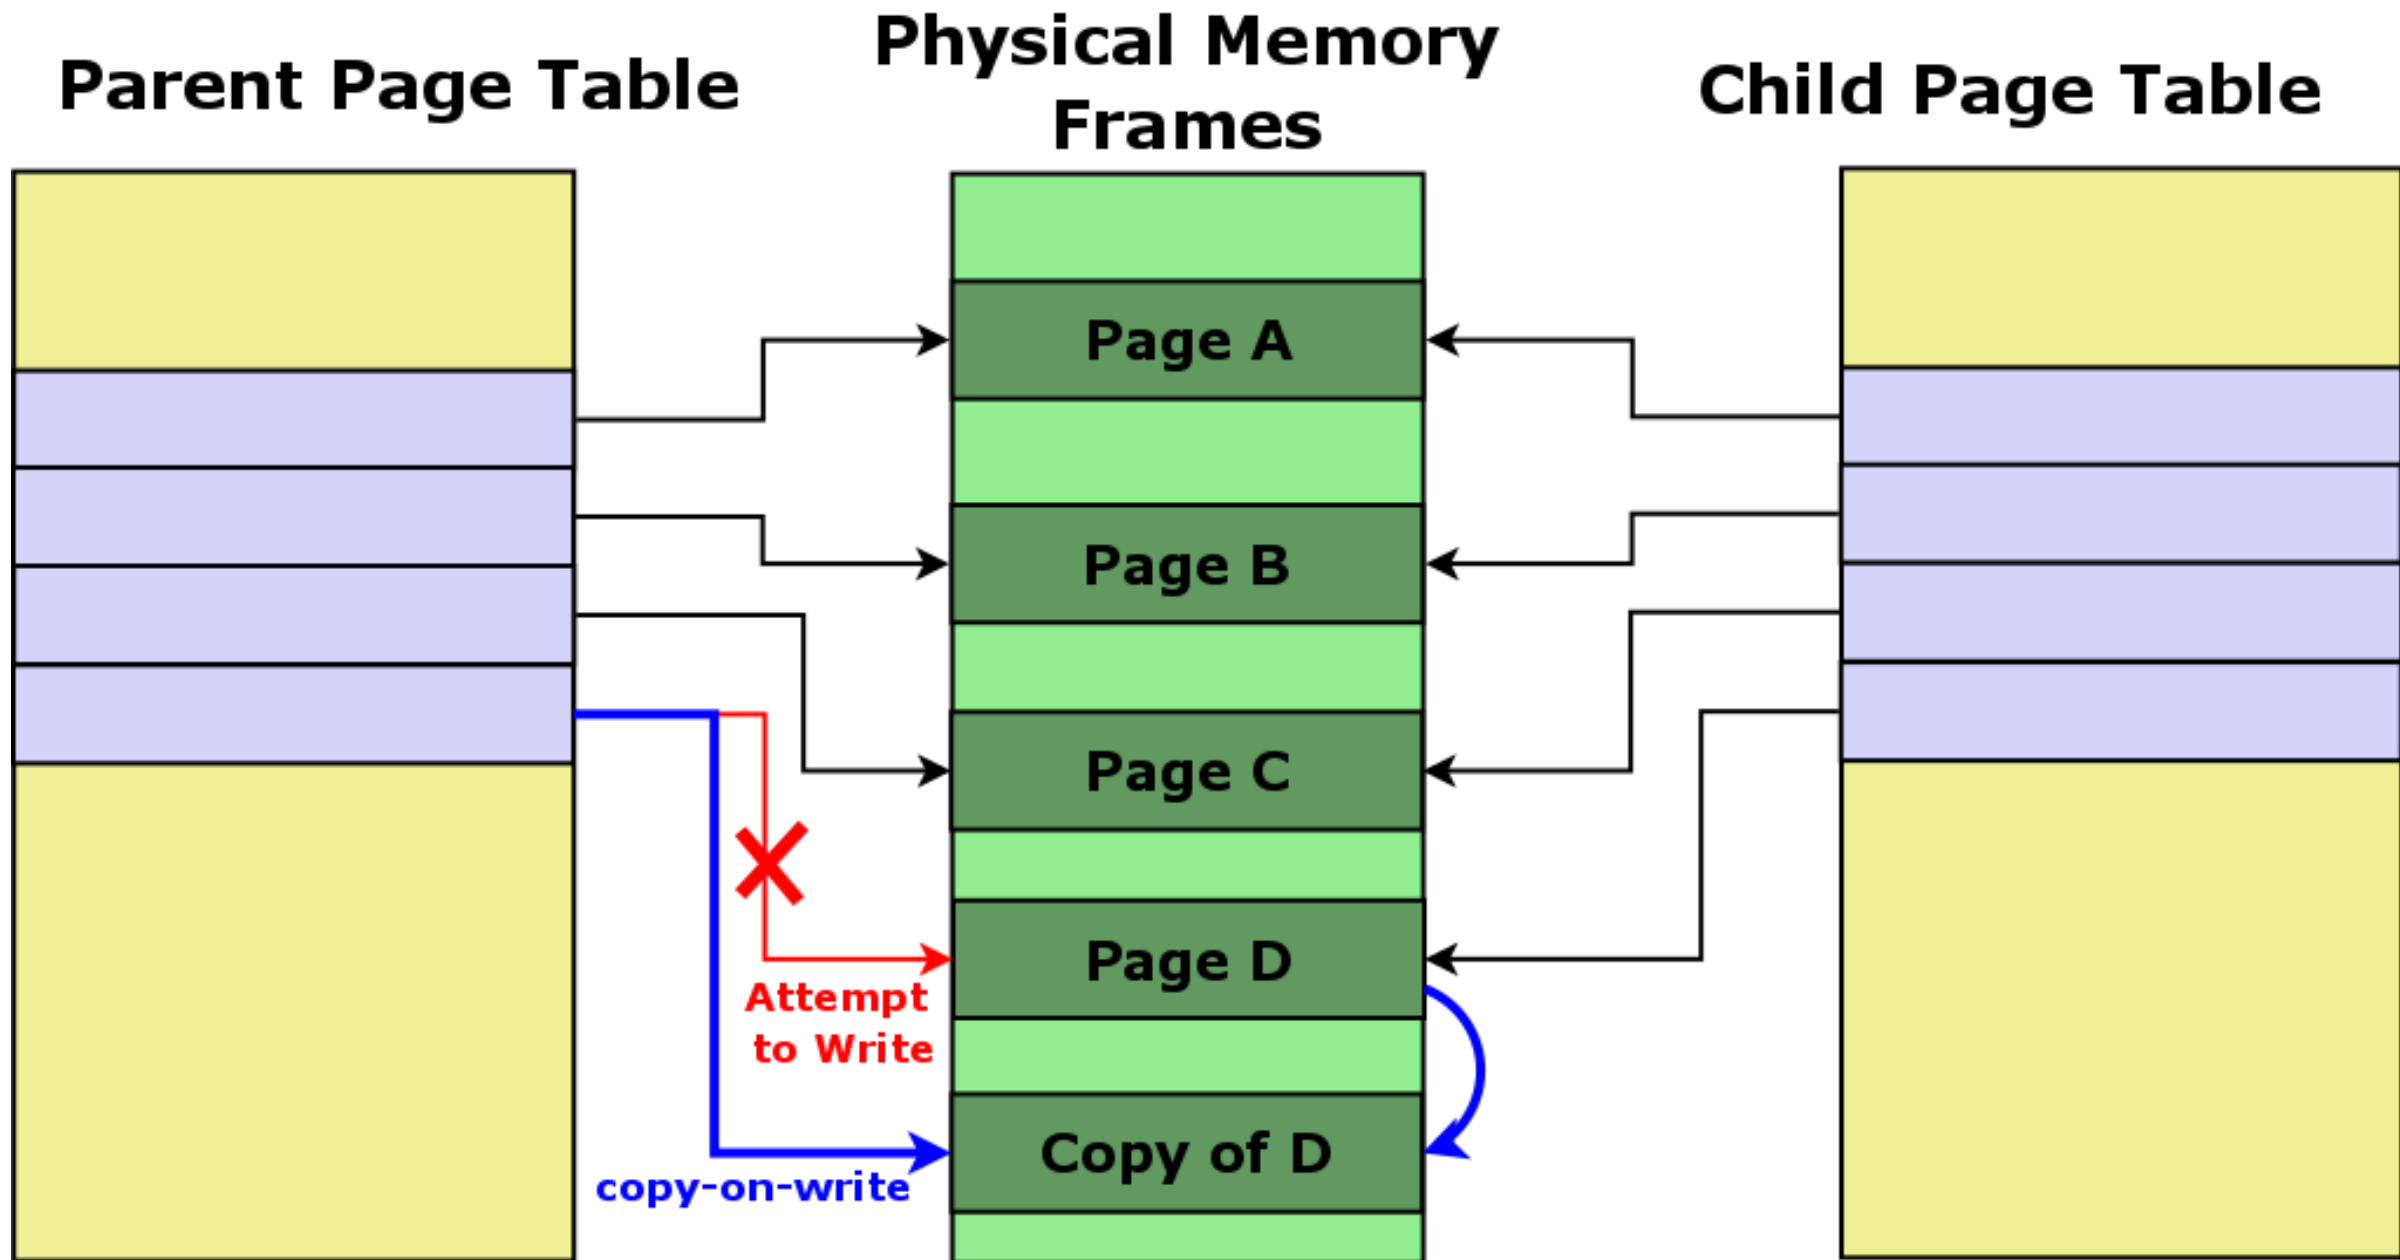


Shared Pages (fork system call → copy-on-write)

- Sharing data is trickier than sharing code, but not impossible.
- In UNIX, after a fork system call the parent and the child are required to share both program text and data.
- Each process will have its own page table pointing initially to the same page frames.
- All pages are marked as read-only in the PTE protection bits.
- As soon as the child or the parent start writing, a trap will be generated due to the read-only protection violation.
- The trap handler will:
 - Make a copy of the page that generated the trap.
 - Set the page entry protection bits to READ/WRITE.
 - Link the page table entry of the process causing the trap to point to the new page.



Shared Pages (fork system call → copy-on-write)



The Shell

```
size_t read_command(char *cmd) {
    if(!fgets(cmd, BUFFER_LEN, stdin)) //get command and put it in line
        return 0; //if user hits CTRL+D break
    size_t length = strlen(cmd); // get command length
    if (cmd[length - 1] == '\n') cmd[length - 1] = '\0'; // clear new line
    return strlen(cmd); // return length of the command read
}

int build_args(char * cmd, char ** argv) {
    char *token; //split command into separate strings
    token = strtok(cmd, " ");
    int i=0;
    while(token!=NULL){// loop for all tokens
        argv[i]=token; // store token
        token = strtok(NULL, " "); // get next token
        i++; // increment number of tokens
    }
    argv[i]=NULL; //set last value to NULL for execvp
    return i; // return number of tokens
}

void set_program_path (char * path, char * bin, char * prog){
    memset (path,0,1024); // intialize buffer
    strcpy(path, bin); //copy /bin/ to file path
    strcat(path, prog); //add program to path
    for(int i=0; i<strlen(path); i++) //delete newline
        if(path[i]=='\n') path[i]='\0';
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>

#define BUFFER_LEN 1024

int main(){
    char line[BUFFER_LEN]; //get command line
    char* argv[100]; //user command
    char* bin= "/bin/"; //set path at bin
    char path[1024]; //full file path
    int argc; //arg count
    while(1){
        printf("My shell>> "); //print shell prompt
        if (read_command(line) == 0 )
            {printf("\n"); break;} // CTRL+D pressed
        if (strcmp(line, "exit") == 0) break; //exit
        argc = build_args (line,argv); // build program argument
        set_program_path (path,bin,argv[0]); // set program full path
        int pid= fork(); //fork child
        if(pid==0){ //Child
            execve(path,argv,0); // if failed process is not replaced
            // then print error message
            fprintf(stderr, "Child process could not do execve\n");
        }else wait(NULL); //Parent
    }
    return 0;
}
```

